

TRAINING GENERATIVE ADVERSARIAL NETWORKS WITH BINARY NEURONS BY END-TO-END BACKPROPAGATION

Hao-Wen Dong and Yi-Hsuan Yang

OUTLINES

- ◆ **Backgrounds**
 - ◆ Generative Adversarial Networks
 - ◆ Binary Neurons
 - ◆ Straight-through Estimators
- ◆ **BinaryGAN**
- ◆ **Experiments & Results**
- ◆ **Discussions & Conclusion**

BACKGROUNDS


GENERATIVE ADVERSARIAL NETWORKS

- ◆ **Goal**—learn a **mapping** from the prior distribution to the data distribution [2]




GENERATIVE ADVERSARIAL NETWORKS

- ◆ **Goal**—learn a **mapping** from the prior distribution to the data distribution [2]



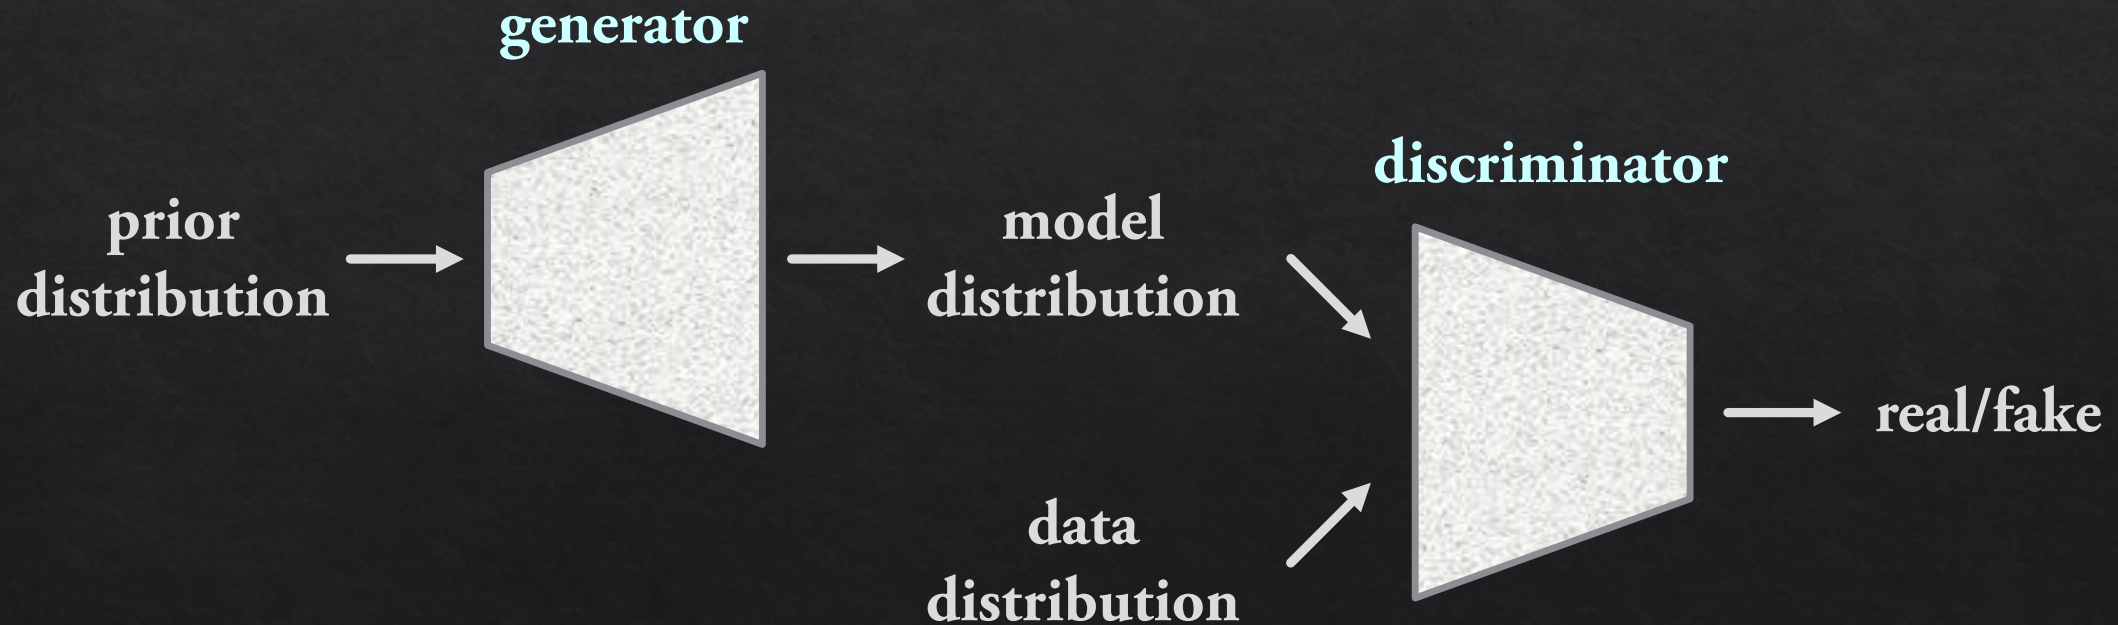
GENERATIVE ADVERSARIAL NETWORKS

- ◆ Use a deep neural network to learn an **implicit** mapping



GENERATIVE ADVERSARIAL NETWORKS

- ◇ Use another deep neural network to provide **guidance/critics**



BINARY NEURONS

◊ **Definition:** neurons that output **binary-valued** predictions

BINARY NEURONS

- ◇ **Definition:** neurons that output **binary-valued** predictions
- ◇ **Deterministic binary neurons (DBNs):**

$$DBN(x) \equiv \begin{cases} 1, & \text{if } \sigma(x) < 0.5 \\ 0, & \text{otherwise} \end{cases} \quad (\text{hard thresholding})$$

BINARY NEURONS

◇ **Definition:** neurons that output **binary-valued** predictions

◇ **Deterministic binary neurons (DBNs):**

$$DBN(x) \equiv \begin{cases} 1, & \text{if } \sigma(x) < 0.5 \\ 0, & \text{otherwise} \end{cases} \quad (\text{hard thresholding})$$

◇ **Stochastic binary neurons (SBNs):**

$$SBN(x) \equiv \begin{cases} 1, & \text{if } z < \sigma(x) \\ 0, & \text{otherwise} \end{cases}, \quad z \sim U[0, 1] \quad (\text{Bernoulli sampling})$$

BACKPROPAGATING THROUGH BINARY NEURONS

- ◇ Backpropagating through binary neurons is **intractable**
 - ◇ For DBNs, it involves the nondifferentiable threshold function
 - ◇ For SBNs, it requires the computation of expected gradients on all possible combinations (*exponential to the number of binary neurons*) of values taken by the binary neurons

BACKPROPAGATING THROUGH BINARY NEURONS

- ◆ Backpropagating through binary neurons is intractable
 - ◆ For DBNs, it involves the nondifferentiable threshold function
 - ◆ For SBNs, it requires the computation of expected gradients on all possible combinations (*exponential to the number of binary neurons*) of values taken by the binary neurons
- ◆ We can introduce **gradient estimators** for the binary neurons
 - ◆ Examples include Straight-through [3,4] (*the one adopted in this work*), REINFORCE [5], REBAR [6], RELAX [7] estimators

STRAIGHT-THROUGH ESTIMATORS

- ◆ **Straight-through estimators:** [3,4]
 - ◆ Forward pass—use hard thresholding (DBNs) or Bernoulli sampling (SBNs)
 - ◆ Backward pass—pretend it as an **identity function**

STRAIGHT-THROUGH ESTIMATORS

- ◆ **Straight-through estimators:** [3,4]
 - ◆ Forward pass—use hard thresholding (DBNs) or Bernoulli sampling (SBNs)
 - ◆ Backward pass—pretend it as an identity function
- ◆ **Sigmoid-adjusted straight-through estimators**
 - ◆ Use **the derivative of the sigmoid function** in the backward pass
 - ◆ Found to achieve better performance in a classification task presented in [4]
 - ◆ Adopted in this work

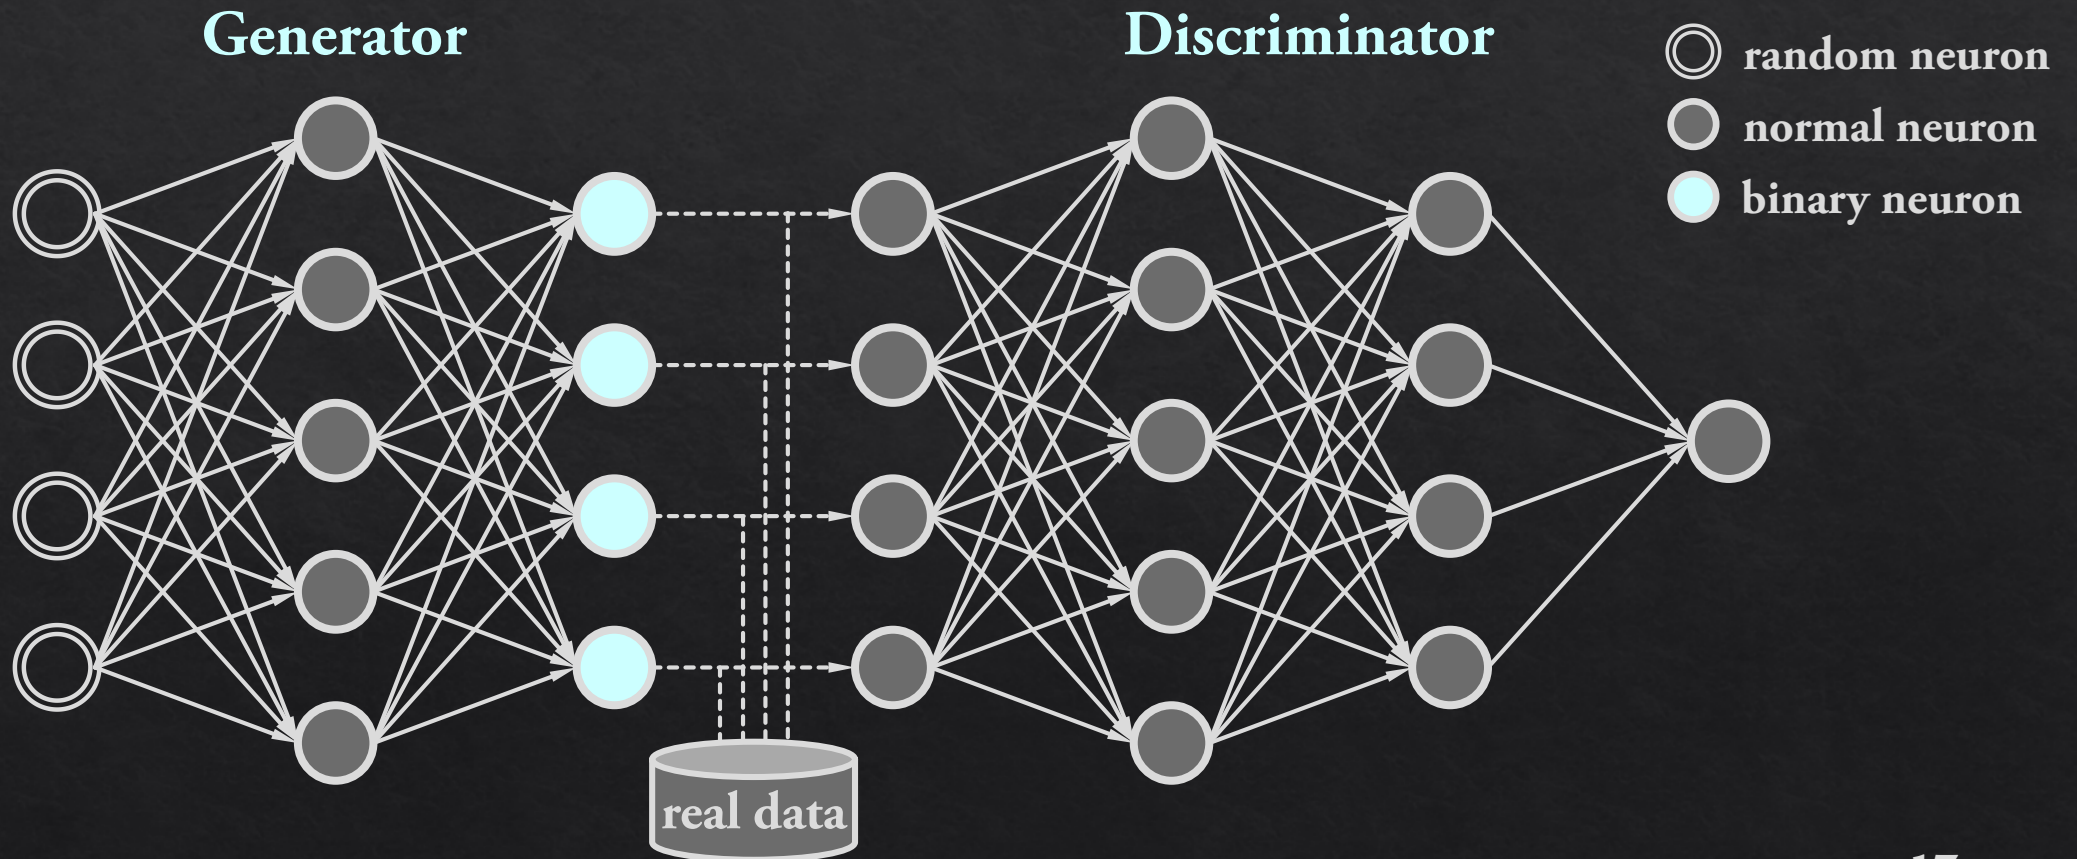
BINARYGAN

BINARYGAN

- ◆ Use binary neurons at the **output layer** of the generator
- ◆ Use sigmoid-adjusted straight-through estimators to provide the gradients for the binary neurons
- ◆ Train the whole network by **end-to-end backpropagation**

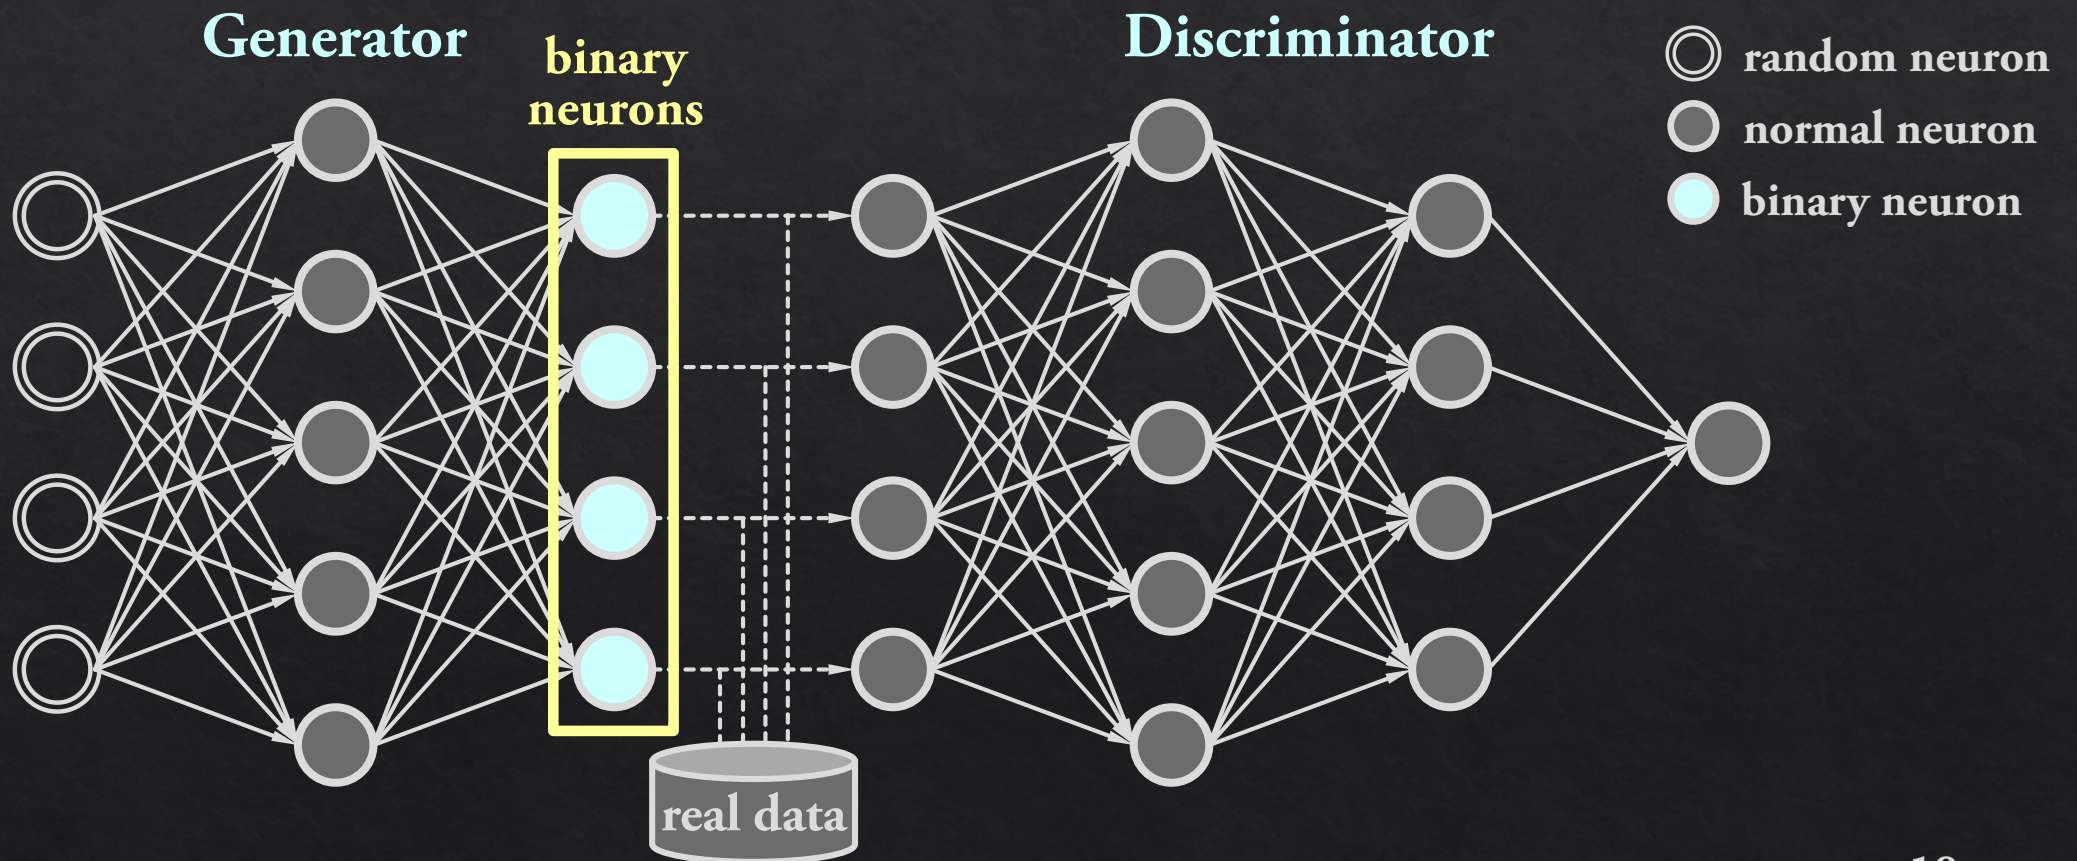
BINARYGAN

(implemented by multilayer perceptrons)



BINARYGAN

(implemented by multilayer perceptrons)




EXPERIMENTS & RESULTS

TRAINING DATA

- ◆ **Binarized MNIST** handwritten digit database [8]
 - ◆ Digits with nonzero intensities \rightarrow 1
 - ◆ Digits with zero intensities \rightarrow 0

Sample binarized MNIST digits



IMPLEMENTATION DETAILS

- ◆ Batch size is 64
- ◆ Use WGAN-GP [9] objectives
- ◆ Use Adam optimizer [10]
- ◆ Apply batch normalization [11] to the generator (but not to the discriminator)
- ◆ Binary neurons are implemented with the code kindly provided in a blog post on the R2RT blog [12]


IMPLEMENTATION DETAILS

- ◆ Apply slope annealing trick [13]
 - ◆ Gradually increase the slopes of the sigmoid functions used in the sigmoid-adjusted straight-through estimators as the training proceeds
 - ◆ We multiply the slopes by 1.1 after each epoch
 - ◆ The slopes start from 1.0 and reach 6.1 at the end of 20 epochs


EXPERIMENT I—DBNs vs SBNs

- ◇ DBNs and SBNs can achieve **similar qualities**
- ◇ They show **distinct characteristics** on the **preactivated outputs**

DBNs




SBNs



EXPERIMENT I—DBNs vs SBNs

Histograms of the preactivated outputs



DBNs

- more values in the middle
- a notch at 0.5 (the threshold)

SBNs

- more values close to 0 and 1


EXPERIMENT II—REAL-VALUED MODEL

- ◇ Use **no binary neurons**
- ◇ Train the discriminator by the **real-valued outputs** of the generator



EXPERIMENT II—REAL-VALUED MODEL

Histograms of the preactivated outputs



DBNs

- more values in the middle
- a notch at 0.5 (the threshold)

SBNs

- more values close to 0 and 1


Real-valued model

- even more U-shaped


EXPERIMENT III—GAN OBJECTIVES

- ◇ WGAN [14] model can achieve similar qualities to the WGAN-GP
- ◇ GAN [2] model suffers from **mode collapse** issue

GAN + DBNs



GAN + SBNs



WGAN + DBNs



WGAN + SBNs




EXPERIMENT IV—MLPs vs CNNs


- ◇ CNN model produces less artifacts even with a small number of trainable parameters (MLP—0.53M; CNN—1.4M)

MLP model

DBNs



SBNs




CNN model

DBNs



SBNs



DISCUSSIONS & CONCLUSION

DISCUSSIONS

- ◆ **Why is binary neurons important?**
 - ◆ Open the possibility of conditional computation graph [4,13]
 - ◆ Move toward a stronger AI that can make reliable decisions

DISCUSSIONS

- ◆ **Why is binary neurons important?**
 - ◆ Open the possibility of conditional computation graph [4,13]
 - ◆ Move toward a stronger AI that can make reliable decisions
- ◆ **Other approaches to model discrete distributions with GANs**
 - ◆ Replace the target discrete outputs with continuous relaxations
 - ◆ View the generator as agent in reinforcement learning (RL) and introduce RL-based training strategies

FUTURE WORK

- ◆ Examine the use of gradient estimators for training a GAN that has a **conditional computation graph**

CONCLUSION

- ◆ **A new GAN model that**
 - ◆ can generate binary-valued predictions without further post-processing
 - ◆ can be trained by end-to-end backpropagation
- ◆ **Experimentally compare**
 - ◆ deterministic and stochastic binary neurons
 - ◆ the proposed model and the real-valued model
 - ◆ GAN, WGAN, WGAN-GP objectives
 - ◆ MLPs and CNNs

REFERENCES

- [1] Hao-Wen Dong and Yi-Hsuan Yang. Training Generative Adversarial Networks with Binary Neurons by End-to-end Backpropagation. *arXiv Preprint arXiv:1810.04714*, 2018
- [2] Ian J. Goodfellow et al. Generative adversarial nets. In *Proc. NIPS*, 2014.
- [3] Geoffrey Hinton. Neural networks for machine learning—Using noise as a regularizer (lecture 9c), 2012. Coursera, video lectures. [Online] <https://www.coursera.org/lecture/neural-networks/using-noise-as-a-regularizer-7-min-wbw7b>.
- [4] Yoshua Bengio, Nicholas Léonard, and Aaron C. Courville. Estimating or propagating gradients through stochastic neurons for conditional computation. *arXiv preprint arXiv:1308.3432*, 2013.
- [5] Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning*, 8(3-4):229–256, 1992.
- [6] George Tucker, Andriy Mnih, Chris J Maddison, and Jascha Sohl-Dickstein. REBAR: Low-variance, unbiased gradient estimates for discrete latent variable models. In *Proc. NIPS*, 2017.
- [7] Will Grathwohl, Dami Choi, Yuhuai Wu, Geoffrey Roeder, and David Duvenaud. Backpropagation through the Void: Optimizing control variates for black-box gradient estimation. In *Proc. ICLR*. 2018.

REFERENCES

- [8] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proc. IEEE*, 86(11):2278–2324, 1998.
- [9] Ishaan Gulrajani, Faruk Ahmed, Martin Arjovsky, Vincent Dumoulin, and Aaron Courville. Improved training of Wasserstein GANs. In *Proc. NIPS*, 2017.
- [10] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [11] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *Proc. ICML*, 2015.
- [12] Binary stochastic neurons in tensorflow, 2016. Blog post on the R2RT blog. [Online] <https://r2rt.com/binary-stochastic-neurons-in-tensorflow>.
- [13] Junyoung Chung, Sungjin Ahn, and Yoshua Bengio. Hierarchical multiscale recurrent neural networks. In *Proc. ICLR*, 2017.
- [14] Martin Arjovsky, Soumith Chintala, and Léon Bottou. Wasserstein generative adversarial networks. In *Proc. ICML*, 2017.

Thank you for your attention